

ASSIGNMENT 1

End-of-Sentence Detection and Text Segment Classification

Due date: Sunday March 1 2020 11:59 PM

Part 1: End-of-Sentence Detection

The goal of the first part of the assignment is to create an algorithm for determining whether a given period (".") in a text indicates an end of sentence or just an abbreviation marker. The following examples illustrate some of the difficulties encountered in this distinction:

NOT-END-OF-SENTENCE	119	l viewpoint , as David	C. Robinson has recently shown , t
NOT-END-OF-SENTENCE	128	evealed , '' by Arthur	C. Clarke , Gentry Lee (Bantam)
NOT-END-OF-SENTENCE	136	E . The group led by	C. Delores Tucker , head of the NA
NOT-END-OF-SENTENCE	147	ence and Electronics ;	C. Scott Kulicke , on behalf of Se
NOT-END-OF-SENTENCE	184	g Committee chaired by	C. Rubbia . The report covers stat
END-OF-SENTENCE	192	occurs at 440 degrees	C. A hydrogenation test was carrie
END-OF-SENTENCE	210	ystem at 25 and 50 deg	C. Isotherms consist of five branc
END-OF-SENTENCE	239	C while not at 40 deg	C. Minima on the S/sub Pu / vs. C/
NOT-END-OF-SENTENCE	247	ellulases . Culture of	C. thermocellum will be optimized
NOT-END-OF-SENTENCE	255	the cellulase genes of	C. cellulolyticum and those from t
END-OF-SENTENCE	258	anging from 200 to 300	C. A system developed by the autho
END-OF-SENTENCE	262	ourse on programming in	C. Finally, those who are interest
END-OF-SENTENCE	300	a house on 2213 Perry Dr.	Then the Thomases were seen in
NOT-END-OF-SENTENCE	330	. <P> Early in 1980	Dr. Thomas B. Reed of SERI and Pro

To help you develop a classifier for this distinction, an example set of 45,000 periods and their surrounding context has been provided, each one labelled as EOS (End-of-sentence) or NEOS (Not-end-of-sentence). The examples were extracted primarily from the Brown Corpus and are located in the sent.data.train on Piazza.

For easy manipulation, the training examples have been divided into tab-delimited columns containing the following information:

- Column 1: EOS or NEOS, indicating whether the period in that line marks an end of sentence marker or not.
- Column 2: The ID number of the sentence.
- Columns 3-9: The ± 3 -word surrounding context of the period.
- Column 10: The number of words to the left of the period before the next *reliable* sentence delimiter (e.g. ?, ! or a paragraph marker <P>).
- Column 11: The number of words to the right of the period before the next *reliable* sentence delimiter (e.g. ?, ! or a paragraph marker <P>).
- Column 12: The number of spaces following the period in the original text.

An example of the first 8 columns for the data above is:

TAG	ID#	-3	-2	-1	0	+1	+2
=====							
NEOS	119	as	David	C	.	Robinson	has
NEOS	128	by	Arthur	C	.	Clarke	,
NEOS	136	led	by	C	.	Delores	Tucker
NEOS	147	electronics	;	C	.	Scott	Kulicke
NEOS	184	chaired	by	C	.	Rubbia	.
EOS	192	440	degrees	C	.	A	hydrogenation
EOS	210	50	deg	C	.	Isotherms	consist
EOS	239	40	deg	C	.	Minima	on
NEOS	247	Culture	of	C	.	thermocellum	will
NEOS	255	genes	of	C	.	cellulolyticum	and
EOS	258	to	300	C	.	A	system
EOS	262	programming	in	C	.	Finally	,
EOS	300	2213	Perry	Dr	.	Then	the
NEOS	330	in	1980	Dr	.	Thomas	B

The classifier you develop should be able to take data of this format and predict whether the correct label is EOS or NEOS.

You must empirically derive a decision procedure from the data using a machine learning algorithm such as a decision tree or neural net.

To test the effectiveness of your program, your code will be applied to another file of 5,000 different examples in identical format (`sent.data.test`).

For maximum fairness of the test, you will not be able to see this test data in advance.

To assist you, we have included other files containing wordlists of abbreviations, titles, unlikely proper nouns, and other terms. You may find these useful when building your classifier. To simplify testing and grading, your code must be runnable using the following command:

```
python hw1a.py --train traindata --test testdata --output outputfile
```

where `traindata` is the path to the training data, `testdata` is the path to the test file, and `outputfile` is a file that will contain the line-by-line classification from your algorithm. We have provided starter code that conforms to this specification and already handles loading the data. In addition, it will compute and print out accuracy on your test data. If your code needs to load any other files, please use relative paths, not absolute paths. We will run your program using Python 3.7.2. If you use any external packages, let us know in your writeup (Part 3) so we can install them if needed. We included two sample classifier implementations. the first one (`EOSClassifier1`) trains a model to always predict the most common label that it has seen in the training data. If you use `sent.train` as both the training and test data, using the following command,

```
python hw1a.py --train sent.train --test sent.train
```

you'll see that the performance on the training data is over 90%. This reveals that the data is heavily unbalanced. The second classifier (`EOSClassifier2`) doesn't look at the training data at all! It simply examines the word immediately to the left of the period. If this word

is an abbreviation in the provided abbreviations wordlist, then it predicts NEOS, otherwise EOS. This classifier gets 95% accuracy on the training data. Nevertheless, there is still room to improve. However, beware of overfitting. Just because your classifier gets 99.9% accuracy on the training data does not mean it will perform similarly on the test data. As mentioned before, we recommend that you split your training data into a training set and development set and only use the development set for testing. You'll notice that the starter code uses scikit-learn. Feel free to use scikit-learn or your favorite ML package. One benefit of scikit-learn is that you can easily try out many different ML models, and we encourage you to do so. The official tutorial can be found at <https://scikit-learn.org/stable/tutorial/basic/tutorial.html>

Part 2: Text Segment Classification

Much of the text encountered in real-world NLP systems is intermixed with non-textual components such as tables, figures, formulae, and email/netnews headers. Text itself may be standard paragraph style prose or specialized textual segments such as headlines or section headers, addresses, quoted text or email signature blocks. It is useful to distinguish these different segments, both for processing in IR or message routing systems and for obtaining clean prose as training data for language models.

The goal of this part of the assignment is to label each line in a text file with the segment type of the text block the line appears in. For example:

```
NNHEAD From: desmedt@ruls40.Berkeley.EDU (Koenraad De Smedt)
NNHEAD Newsgroups: comp.ai.nat-lang
NNHEAD Subject: CFP: 5th European Workshop on Natural Language Generation
NNHEAD Date: 24 Oct 1994 09:36:40 GMT
NNHEAD Organization: Leiden University
NNHEAD Message-ID: <38fv78$9du@highway.LeidenUniv.nl>
NNHEAD Keywords: Natural Language Generation Workshop
#BLANK#
HEADL                                CALL FOR PAPERS
#BLANK#
HEADL          5th European Workshop on Natural Language Generation
#BLANK#
HEADL                                20-23 May 1995
HEADL                                Leiden, The Netherlands
#BLANK#
PTEXT This workshop aims to bring together researchers interested in Natural
PTEXT Language Generation from such different perspectives as linguistics,
PTEXT artificial intelligence, psychology, and engineering. The meeting
PTEXT continues the tradition of a series of workshops held biannually in
PTEXT Europe (Royaumont, 1987; Edinburgh, 1989; Judenstein, 1991 and Pisa,
PTEXT 1993) but open to researchers from all over the world.
#BLANK#
HEADL                                Programme
#BLANK#
PTEXT Papers, posters and demonstrations are invited on original and
PTEXT substantial work related to the automatic generation of natural
PTEXT language, including computer linguistics research, artificial
PTEXT intelligence methods, computer models of human language processing,
#BLANK#
```

```

PTEXT  All contributions should be sent BEFORE 1 JANUARY 1995 to the
PTEXT  Programme Chairman at the following address:
#BLANK#
ADDRESS      Philippe Blache
ADDRESS      2LC - CNRS
ADDRESS      1361 route des Lucioles
ADDRESS      F-06560 Sophia Antipolis
ADDRESS      tel : +33 92.96.73.98
ADDRESS      fax : +33 93.65.29.27
ADDRESS      e-mail : pb@llaor.unice.fr
#BLANK#
#BLANK#
QUOTED > SJC (San Jose, CA) has an open observation deck on its older terminal
QUOTED > A. I have not used this terminal in quite some time, so I don't know
QUOTED > if sightseers still have access to it. The problem with this deck was
QUOTED > that the @#$%^& restaurant would block your view just as the jets
QUOTED > were getting off the ground. Nevertheless, you could still watch the
QUOTED > planes taxi and then accelerate from the start of the runway.
QUOTED >
QUOTED > Why is it that the newer terminals no longer have these outdoor viewing
QUOTED > areas? Security, I suppose. A sad sign of our times.
QUOTED >
#BLANK#
#BLANK#
SIG      '',''
SIG      (0 0)
SIG      +-----00o--(_)--o00-----+
SIG      |                               |
SIG      | Rajeev Agarwal                | "Hanging in there..." |
SIG      |                               |
SIG      | Dept. of Computer Science      | (601) 325-8073 or 2756 (Off.) |
SIG      | Mississippi State University  | (601) 325-7506 (Lab)   |
SIG      +-----+-----+
SIG      (_) (_)
```

A description of the different segment types and examples of them may be found in the Homework directory in the segments subdirectory. Many segments may be classified by the presence of relatively simple patterns within the segment, such as Message-ID: or From: in a netnews header NNHEAD, though other segments may be harder to distinguish. e standards for classification are located in the "standard" file, but dont worry about conforming too closely. Priority will be placed on the creativity and completeness of the segment classifier, not on conforming to any arbitrary definitions of what constitutes a signature or table, for example. We have provided starter code that trains a decision tree classifier. We preprocess each input into a feature vector of four features:

- number of characters
- number of characters after trimming whitespace
- number of words
- 1 or 0 depending if a > is present

With only these four features, this classifier gets 85% accuracy on the training data. This method of manual feature extraction is one possible way to approach this problem.

You may also consider using an algorithm to automatically learn features. Similar to Part 1, we must be able to run your code using the following commands:

```
python hw1b.py --train traindata --test testdata --output outputfile --format line
python hw1b.py --train traindata --test testdata --output outputfile --format segment
```

where the `-output` switch indicates whether your program should classify by line or by segment.

Part 3: Writeup

Create a short writeup documenting your algorithms and approach for Parts 1 and 2. Describe the models you used, report their performance on the training data, and comment on anything else you found interesting (what did/didn't work, other methods you tried, etc.).

Evaluation:

Submissions will be evaluated as follows:

- 30% - PART 1: Quality, completeness and creativity of algorithm
- 5% - PART 1: Performance on training data
- 20% - PART 1: Performance on independent test data
- 20% - PART 2: Quality, completeness and creativity of algorithm
- 5% - PART 2: Performance on training data
- 15% - PART 2: Performance on independent test data
- 5% - PART 3: Writeup

Submission

We are using Gradescope for submitting assignments. Please submit a compressed file of a folder containing the files for this assignment. Please name follow the naming convention **JHUID_HW1** for the directory. This file should include **all files** that we provide (training data, vocabulary files, etc) as well as your `hw1a.py` and `hw1b.py` files.